

# Predicting NHL players' salaries

UH

3/19/2021

```
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(gridExtra)

##
## Attaching package: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine

library(corrplot)

## corrplot 0.84 loaded

library(randomForest)

## randomForest 4.6-14

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:gridExtra':
##
##   combine
```

```

## The following object is masked from 'package:dplyr':
##
##   combine

## The following object is masked from 'package:ggplot2':
##
##   margin

# Library(extraTrees)
# Library(rJava)
library(gbm)

## Loaded gbm 2.1.8

library(xgboost)

##
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##   slice

library(h2o)

##
## -----
##
## Your next step is to start H2O:
##   > h2o.init()
##
## For H2O package documentation, ask for help:
##   > ??h2o
##
## After starting H2O, you can use the Web UI at http://localhost:54321
## For more information visit https://docs.h2o.ai
## -----
##
## Attaching package: 'h2o'

## The following objects are masked from 'package:stats':
##
##   cor, sd, var

## The following objects are masked from 'package:base':
##
##   %*%, %in%, &&, ||, apply, as.factor, as.numeric, colnames,
##   colnames<-, ifelse, is.character, is.factor, is.numeric, log,
##   log10, log1p, log2, round, signif, trunc

```

```
setwd("~/Datascience/Data Science Courses/HarvardX Course/Individual project"
)
```

```
# Data for the 2020/21 season from HockeyReference (https://www.hockey-reference.com/friv/current_nhl_salaries.cgi) and NHL (http://www.nhl.com/stats/skaters).
```

```
dat = read.csv("NHL_players_stats_merged.csv", header = T)
```

```
head(dat)
```

```
##      i..Player Team   Salary S.C Pos GP  G  A  P plus_minus PIM PperGP
EVG
## 1 Auston Matthews  TOR 15900000  L  C 27 21 15 36          8   6   1.33
13
## 2 Mitchell Marner  TOR 15000000  R  R 30 11 28 39          15  14   1.30
11
## 3 Connor McDavid  EDM 14000000  L  C 30 17 35 52          10  14   1.73
11
## 4 Artemi Panarin   NYR 13000000  R  L 15  5 14 19           0   0   1.27
3
## 5 Mikko Rantanen   COL 12000000  L  R 26 13 15 28           7   6   1.08
8
## 6 Erik Karlsson    SJS 12000000  R  D 21  2  7  9           1   4   0.43
0
##      EVP PPG PPP SHG SHP OTG GWG  S  TOI
## 1  26   8  10   0   0   1   7 112 1323
## 2  25   0  13   0   1   0   2  82 1358
## 3  30   6  22   0   0   1   4 119 1349
## 4  12   2   7   0   0   0   0  45 1229
## 5  18   5  10   0   0   0   3  87 1210
## 6   5   2   4   0   0   0   0  41 1452
```

```
# Legend: S.C=Skater shoots, Pos=Player position, GP=Games played, G=Goals, A=Assists, P=Points, plus_minus, PIM=Penalty minutes, PperGP=Points per game played, EVG=Even strength goals, EVP=Even strength points, PPG=Powerplay goals, PPP=Powerplay points, SHG=Shorthanded goals, SHP=Short handed points, OTG=Overtime goals, GWG=Game-winning goals, S=Shots, TOI=Time on ice
```

## Introduction

The goal of this project is to predict the salaries of NHL players. I am using data from the current 2020/21 season, which is publicly available. I am new to machine learning and am exploring various methods and approaches. I am a hockey fan, hence the topic. The csv files are posted on github.

Legend: S.C=Skater shoots, Pos=Player position, GP=Games played, G=Goals, A=Assists, P=Points, plus\_minus, PIM=Penalty minutes, PperGP=Points per game played, EVG=Even strength goals, EVP=Even strength points, PPG=Powerplay goals, PPP=Powerplay points, SHG=Shorthanded goals, SHP=Shorthanded points, OTG=Overtime goals, GWG=Game-winning goals, S=Shots, TOI=Time on ice

```
# Wrangling and getting descriptive numbers
is.na(dat$salary)

## logical(0)

cat("\n\n Sort data frame by salary in descending order\n")

##
##
## Sort data frame by salary in descending order

# sort data frame by salary in descending order
dat_sorted <- dat[with(dat, order(-dat$Salary)), ]

# print(dat_sorted)

# Make sure numbers are in numeric format
dat_sorted$Salary <- as.numeric(dat_sorted$Salary)
dat_sorted$GP <- as.numeric(dat_sorted$GP)
dat_sorted$G <- as.numeric(dat_sorted$G)
dat_sorted$A <- as.numeric(dat_sorted$A)
dat_sorted$plus_minus <- as.numeric(dat_sorted$plus_minus)
dat_sorted$PIM <- as.numeric(dat_sorted$PIM)
dat_sorted$PperGP <- as.numeric(dat_sorted$PperGP)
dat_sorted$EVG <- as.numeric(dat_sorted$EVG)
dat_sorted$EVP <- as.numeric(dat_sorted$EVP)
dat_sorted$PPG <- as.numeric(dat_sorted$PPG)
dat_sorted$PPP <- as.numeric(dat_sorted$PPP)
dat_sorted$SHG <- as.numeric(dat_sorted$SHG)
dat_sorted$SHP <- as.numeric(dat_sorted$SHP)
dat_sorted$OTG <- as.numeric(dat_sorted$OTG)
dat_sorted$GWG <- as.numeric(dat_sorted$GWG)
dat_sorted$S <- as.numeric(dat_sorted$S)
dat_sorted$TOI <- as.numeric(dat_sorted$TOI)

# Examine the structure of the dat dataset
str(dat)
```

```
## 'data.frame': 593 obs. of 22 variables:
## $ i..Player : chr "Auston Matthews" "Mitchell Marner" "Connor McDavid" "
Artemi Panarin" ...
## $ Team : chr "TOR" "TOR" "EDM" "NYR" ...
## $ Salary : int 15900000 15000000 14000000 13000000 12000000 12000000
12000000 12000000 10570000 10570000 ...
## $ S.C : chr "L" "R" "L" "R" ...
## $ Pos : chr "C" "R" "C" "L" ...
## $ GP : int 27 30 30 15 26 21 30 27 27 1 ...
## $ G : int 21 11 17 5 13 2 9 10 10 0 ...
## $ A : int 15 28 35 14 15 7 17 20 15 1 ...
## $ P : int 36 39 52 19 28 9 26 30 25 1 ...
## $ plus_minus: int 8 15 10 0 7 1 11 -1 7 0 ...
## $ PIM : int 6 14 14 0 6 4 6 4 12 0 ...
## $ PperGP : num 1.33 1.3 1.73 1.27 1.08 0.43 0.87 1.11 0.93 1 ...
## $ EVG : int 13 11 11 3 8 0 4 7 6 0 ...
## $ EVP : int 26 25 30 12 18 5 17 19 16 1 ...
## $ PPG : int 8 0 6 2 5 2 5 3 3 0 ...
## $ PPP : int 10 13 22 7 10 4 9 11 8 0 ...
## $ SHG : int 0 0 0 0 0 0 0 0 1 0 ...
## $ SHP : int 0 1 0 0 0 0 0 0 1 0 ...
## $ OTG : int 1 0 1 0 0 0 0 0 0 0 ...
## $ GWG : int 7 2 4 0 3 0 1 0 2 0 ...
## $ S : int 112 82 119 45 87 41 83 49 66 1 ...
## $ TOI : int 1323 1358 1349 1229 1210 1452 1090 1132 1160 638 ...
```

*# Create a summary for the dat dataset*  
summary(dat)

```
## i..Player      Team      Salary      S.C
## Length:593     Length:593   Min.   : 700000   Length:593
## Class :character Class :character 1st Qu.: 925000   Class :character
## Mode  :character Mode  :character Median : 2050000   Mode  :character
##                                     Mean  : 2928660
##                                     3rd Qu.: 4100000
##                                     Max.   :15900000
##      Pos      GP      G      A
## Length:593    Min.   : 1.00   Min.   : 0.000   Min.   : 0.00
## Class :character 1st Qu.:19.00   1st Qu.: 1.000   1st Qu.: 2.00
## Mode  :character Median :24.00   Median : 3.000   Median : 5.00
##                                     Mean  :22.28   Mean  : 3.826   Mean  : 6.41
##                                     3rd Qu.:27.00   3rd Qu.: 6.000   3rd Qu.:10.00
##                                     Max.   :31.00   Max.   :21.000   Max.   :35.00
##      P      plus_minus      PIM      PperGP
## Min.   : 0.00   Min.   : -24.0000   Min.   : 0.000   Min.   :0.0000
## 1st Qu.: 4.00   1st Qu.: -4.0000   1st Qu.: 4.000   1st Qu.:0.1900
## Median : 8.00   Median :  0.0000   Median : 8.000   Median :0.3600
## Mean  :10.24   Mean  :  0.1046   Mean  : 9.997   Mean  :0.4245
## 3rd Qu.:15.00   3rd Qu.:  3.0000   3rd Qu.:14.000   3rd Qu.:0.6100
## Max.   :52.00   Max.   : 25.0000   Max.   :53.000   Max.   :1.7300
```

```
##           EVG           EVP           PPG           PPP
## Min.      : 0.000   Min.      : 0.000   Min.      : 0.0000   Min.      : 0.000
## 1st Qu.: 1.000   1st Qu.: 3.000   1st Qu.: 0.0000   1st Qu.: 0.000
## Median : 2.000   Median : 7.000   Median : 0.0000   Median : 1.000
## Mean      : 2.877   Mean      : 7.567   Mean      : 0.8583   Mean      : 2.497
## 3rd Qu.: 4.000   3rd Qu.:11.000   3rd Qu.: 1.0000   3rd Qu.: 4.000
## Max.      :13.000   Max.      :30.000   Max.      :10.0000   Max.      :22.000
##           SHG           SHP           OTG           GWG
## Min.      :0.00000   Min.      :0.000   Min.      :0.00000   Min.      :0.0000
## 1st Qu.:0.00000   1st Qu.:0.000   1st Qu.:0.00000   1st Qu.:0.0000
## Median :0.00000   Median :0.000   Median :0.00000   Median :0.0000
## Mean      :0.09106   Mean      :0.172   Mean      :0.09275   Mean      :0.6054
## 3rd Qu.:0.00000   3rd Qu.:0.000   3rd Qu.:0.00000   3rd Qu.:1.0000
## Max.      :4.00000   Max.      :4.000   Max.      :3.00000   Max.      :7.0000
##           S           TOI
## Min.      : 0.00   Min.      : 248.0
## 1st Qu.: 21.00   1st Qu.: 796.0
## Median : 35.00   Median : 973.0
## Mean      : 38.29   Mean      : 973.8
## 3rd Qu.: 53.00   3rd Qu.:1132.0
## Max.      :132.00   Max.      :1614.0
```

Now we have some basic numbers about the dataset. There are 593 skaters for which I could match stats from the NHL.com website and salaries. The highest salaries are for Matthews, Marner, and McDavid; these are all forwards. The highest salaries for defensemen are for Karlsson, Aho, and Trouba.

I thought that forwards and defensemen will be evaluated based on different parameters. For instance, goals per season will likely not be the main criterium for a defenseman. Hence I am splitting the dataset into forwards and defensemen. How do they differ in terms of salary? I thought initially that forwards would get paid better on average.

*# Separate forwards and defenders; first forwards*

```
Forwards<-subset(dat, dat$Pos !='D')
head(Forwards)
```

```
##           i..Player Team   Salary S.C Pos GP  G  A  P plus_minus PIM PperGP
EVG
## 1 Auston Matthews   TOR 15900000  L   C 27 21 15 36           8   6   1.33
13
## 2 Mitchell Marner   TOR 15000000  R   R 30 11 28 39          15  14   1.30
11
## 3 Connor McDavid    EDM 14000000  L   C 30 17 35 52          10  14   1.73
11
## 4 Artemi Panarin     NYR 13000000  R   L 15  5 14 19           0   0   1.27
3
## 5 Mikko Rantanen     COL 12000000  L   R 26 13 15 28           7   6   1.08
8
## 7 John Tavares       TOR 12000000  L   C 30  9 17 26          11   6   0.87
4
```

```
##      EVP PPG PPP SHG SHP OTG GWG   S   TOI
## 1    26   8  10   0   0   1   7 112 1323
## 2    25   0  13   0   1   0   2  82 1358
## 3    30   6  22   0   0   1   4 119 1349
## 4    12   2   7   0   0   0   0  45 1229
## 5    18   5  10   0   0   0   3  87 1210
## 7    17   5   9   0   0   0   1  83 1090
```

*# Now create a dataframe with the defensemen*



```
Def<-subset(dat, dat$Pos == 'D')
head(Def)
```

```
##      i..Player Team   Salary S.C Pos GP G  A  P plus_minus PIM PperGP E
VG
## 6      Erik Karlsson SJS 12000000  R  D 21 2  7  9          1   4   0.43
0
## 10     Sebastian Aho  NYI 10570000  L  D  1 0  1  1          0   0   1.00
0
## 11      Jacob Trouba  NYR 10000000  R  D 18 0  5  5         -1  14   0.28
0
## 14      Drew Doughty  LAK 10000000  R  D 27 6 16 22          0  12   0.81
2
## 15      Brent Burns  SJS 10000000  R  D 25 5  9 14        -13  16   0.56
3
## 24     Jared Spurgeon  MIN  9000000  R  D 25 0  5  5          1   0   0.20
0
##      EVP PPG PPP SHG SHP OTG GWG   S   TOI
## 6        5   2   4   0   0   0   0 41 1452
## 10       1   0   0   0   0   0   0  1  638
## 11       5   0   0   0   0   0   0 30 1307
## 14       7   4  14   0   1   0   0 46 1589
## 15       9   2   5   0   0   0   1 67 1614
## 24       5   0   0   0   0   0   0 50 1309
```

First I am getting some basic numbers from the two datasets. I will do some data wrangling as described below. I will then make some graphs and tables to get a better understanding of datasets.

*# Compare the salary in the Forwards and Defensemen datasets with a boxplot*


```
p1 = dat$Salary[which(dat$Pos != 'D')]/1000000
p2 = dat$Salary[which(dat$Pos == 'D')]/1000000
par(mfrow=c(1,2))
boxplot(p1)
boxplot(p2)
```



```
par(mfrow=c(1,1))
boxplot(p1,p2, main = "Forwards vs defensemen", ylab = "Salary in USD (millions)", names = c("Forwards", "Defensemen"), col = "#69b3a2")
```




## Forwards vs defensemen



```
# Make salary histogram
Histo <- ggplot(dat, aes(x=Salary)) +
  geom_histogram(fill= "#69b3a2", col = "black")
Histo

## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



*# Compare the salary in the Forwards and Defensemen datasets with numbers*  
summary(Forwards)

```
##   i..Player      Team      Salary      S.C
## Length:389      Length:389      Min.   : 700000      Length:389
## Class :character Class :character 1st Qu.: 900000      Class :character
## Mode  :character Mode  :character Median : 2150000      Mode  :character
##                                     Mean  : 2975710
##                                     3rd Qu.: 4250000
##                                     Max.   :15900000
##      Pos      GP      G      A
## Length:389      Min.   : 1.00      Min.   : 0.000      Min.   : 0.000
## Class :character 1st Qu.:20.00      1st Qu.: 2.000      1st Qu.: 3.000
## Mode  :character Median :25.00      Median : 4.000      Median : 5.000
##                                     Mean  :22.82      Mean  : 5.064      Mean  : 6.817
##                                     3rd Qu.:27.00      3rd Qu.: 7.000      3rd Qu.:10.000
##                                     Max.   :31.00      Max.   :21.000      Max.   :35.000
##      P      plus_minus      PIM      PperGP
## Min.   : 0.00      Min.   : -17.000000      Min.   : 0.000      Min.   :0.0000
## 1st Qu.: 5.00      1st Qu.: -4.000000      1st Qu.: 4.000      1st Qu.:0.2500
## Median :10.00      Median : 0.000000      Median : 8.000      Median :0.4300
## Mean   :11.88      Mean   : 0.005141      Mean   : 9.715      Mean   :0.4882
```

```
## 3rd Qu.:18.00 3rd Qu.: 3.000000 3rd Qu.:13.000 3rd Qu.:0.7000
## Max. :52.00 Max. : 20.000000 Max. :53.000 Max. :1.7300
## EVG EVP PPG PPP
## Min. : 0.000 Min. : 0.000 Min. : 0.000 Min. : 0.000
## 1st Qu.: 2.000 1st Qu.: 4.000 1st Qu.: 0.000 1st Qu.: 0.000
## Median : 3.000 Median : 8.000 Median : 0.000 Median : 1.000
## Mean : 3.784 Mean : 8.787 Mean : 1.152 Mean : 2.889
## 3rd Qu.: 5.000 3rd Qu.:12.000 3rd Qu.: 2.000 3rd Qu.: 4.000
## Max. :13.000 Max. :30.000 Max. :10.000 Max. :22.000
## SHG SHP OTG GWG
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :0.0000 Median :0.0000 Median :0.0000 Median :0.0000
## Mean :0.1285 Mean :0.2057 Mean :0.1337 Mean :0.8175
## 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:0.0000 3rd Qu.:1.0000
## Max. :4.0000 Max. :4.0000 Max. :3.0000 Max. :7.0000
## S TOI
## Min. : 0.00 Min. : 248.0
## 1st Qu.: 24.00 1st Qu.: 751.0
## Median : 38.00 Median : 891.0
## Mean : 41.74 Mean : 892.8
## 3rd Qu.: 58.00 3rd Qu.:1044.0
## Max. :132.00 Max. :1358.0
```

```
mean(Forwards$Salary)
```

```
## [1] 2975710
```

```
median(Forwards$Salary)
```

```
## [1] 2150000
```

```
max(Forwards$Salary)
```

```
## [1] 15900000
```

```
min(Forwards$Salary)
```

```
## [1] 700000
```

```
summary(Def)
```

```
## i..Player Team Salary S.C
## Length:204 Length:204 Min. : 700000 Length:204
## Class :character Class :character 1st Qu.: 925000 Class :character
## Mode :character Mode :character Median : 2000000 Mode :character
## Mean : 2838944
## 3rd Qu.: 4025044
## Max. :12000000
## Pos GP G A
## Length:204 Min. : 1.00 Min. : 0.000 Min. : 0.000
## Class :character 1st Qu.:17.00 1st Qu.: 0.000 1st Qu.: 2.000
```

```

## Mode :character      Median :24.00      Median : 1.000      Median : 4.000
##                      Mean :21.25      Mean : 1.466      Mean : 5.632
##                      3rd Qu.:27.00      3rd Qu.: 2.000      3rd Qu.: 9.000
##                      Max. :31.00      Max. :11.000      Max. :22.000
##      P      plus_minus      PIM      PperGP
## Min. : 0.000      Min. : -24.0000      Min. : 0.00      Min. :0.0000
## 1st Qu.: 2.000      1st Qu.: -3.0000      1st Qu.: 4.00      1st Qu.:0.1375
## Median : 5.500      Median : 0.0000      Median : 8.00      Median :0.2500
## Mean : 7.098      Mean : 0.2941      Mean :10.53      Mean :0.3032
## 3rd Qu.:10.250      3rd Qu.: 4.0000      3rd Qu.:14.25      3rd Qu.:0.4300
## Max. :26.000      Max. : 25.0000      Max. :45.00      Max. :1.0000
##      EVG      EVP      PPG      PPP
## Min. :0.000      Min. : 0.00      Min. :0.0000      Min. : 0.00
## 1st Qu.:0.000      1st Qu.: 2.00      1st Qu.:0.0000      1st Qu.: 0.00
## Median :1.000      Median : 5.00      Median :0.0000      Median : 0.00
## Mean :1.147      Mean : 5.24      Mean :0.299      Mean : 1.75
## 3rd Qu.:2.000      3rd Qu.: 8.00      3rd Qu.:0.0000      3rd Qu.: 3.00
## Max. :8.000      Max. :16.00      Max. :5.000      Max. :16.00
##      SHG      SHP      OTG      GWG
## Min. :0.00000      Min. :0.0000      Min. :0.00000      Min. :0.000
## 1st Qu.:0.00000      1st Qu.:0.0000      1st Qu.:0.00000      1st Qu.:0.000
## Median :0.00000      Median :0.0000      Median :0.00000      Median :0.000
## Mean :0.01961      Mean :0.1078      Mean :0.01471      Mean :0.201
## 3rd Qu.:0.00000      3rd Qu.:0.0000      3rd Qu.:0.00000      3rd Qu.:0.000
## Max. :1.00000      Max. :2.0000      Max. :1.00000      Max. :2.000
##      S      TOI
## Min. : 0.00      Min. : 568.0
## 1st Qu.:15.75      1st Qu.: 961.5
## Median :30.00      Median :1132.5
## Mean :31.73      Mean :1128.2
## 3rd Qu.:46.25      3rd Qu.:1321.2
## Max. :88.00      Max. :1614.0

mean(Def$Salary)

## [1] 2838944

median(Def$Salary)

## [1] 2e+06

max(Def$Salary)

## [1] 12000000

min(Def$Salary)

## [1] 700000

```

While the top salaries are reserved for forwards, defensemen have similar salaries in terms of mean and median.

The boxplots and histogram tell us a few things: (1) There are many outliers for extremely well paid skaters, and (2) the salaries are not normally distributed. Still, let's try multiple linear regression as a first approach.

```
# Build correlation matrix of all parameters. First exclude non-numeric data.
# For forwards
```

```
sapply(Forwards, is.numeric)
```

```
## i..Player      Team      Salary      S.C      Pos      GP
G
##      FALSE      FALSE      TRUE      FALSE      FALSE      TRUE      TR
UE
##      A      P plus_minus      PIM      PperGP      EVG      E
VP
##      TRUE      TRUE      TRUE      TRUE      TRUE      TRUE      TR
UE
##      PPG      PPP      SHG      SHP      OTG      GWG
S
##      TRUE      TRUE      TRUE      TRUE      TRUE      TRUE      TR
UE
##      TOI
##      TRUE
```


```
Forwards_num_data <- Forwards[, sapply(Forwards, is.numeric)]
cor(Forwards_num_data, use = "complete.obs", method = "pearson")
```

```
##      Salary      GP      G      A      P plus_minu
s
## Salary      1.00000000 0.29855469 0.5688157 0.65067772 0.6755561 0.21067993
8
## GP      0.29855469 1.00000000 0.5412164 0.48593962 0.5565427 0.07406654
1
## G      0.56881565 0.54121637 1.00000000 0.65818317 0.8731069 0.35628175
1
## A      0.65067772 0.48593962 0.6581832 1.00000000 0.9417041 0.38327364
3
## P      0.67555606 0.55654272 0.8731069 0.94170415 1.0000000 0.40741440
7
## plus_minus 0.21067994 0.07406654 0.3562818 0.38327364 0.4074144 1.00000000
0
## PIM      0.09336739 0.33337439 0.1644578 0.14895819 0.1699550 0.00927598
3
## PperGP      0.67497667 0.38663889 0.8178614 0.90653432 0.9525366 0.43419902
6
## EVG      0.45292699 0.53644744 0.9031630 0.56735475 0.7710139 0.40270019
4
## EVP      0.59459280 0.58534638 0.8360137 0.87041134 0.9372564 0.49063225
4
## PPG      0.52182429 0.31318935 0.7397810 0.54454107 0.6832272 0.12262101
2
## PPP      0.63118408 0.36409513 0.7103241 0.82328113 0.8505672 0.16895420
```

3						
## SHG	0.06356257	0.13312371	0.1513386	0.03688819	0.0915189	0.11105985
8						
## SHP	0.08757546	0.16823725	0.1468328	0.08361223	0.1197625	0.15834274
4						
## OTG	0.22738305	0.13913596	0.3497504	0.24571196	0.3154147	0.20458545
8						
## GWG	0.41131968	0.31020592	0.6458213	0.45052559	0.5803562	0.35110797
3						
## S	0.58010268	0.65496266	0.7934833	0.70747705	0.8127386	0.20219035
8						
## TOI	0.67030509	0.53554053	0.7179934	0.77467838	0.8225208	0.21497405
4						
##	PIM	PperGP	EVG	EVP	PPG	
## Salary	0.0933673853	0.67497667	0.45292699	0.59459280	0.521824293	
## GP	0.3333743868	0.38663889	0.53644744	0.58534638	0.313189348	
## G	0.1644578306	0.81786144	0.90316301	0.83601370	0.739781046	
## A	0.1489581870	0.90653432	0.56735475	0.87041134	0.544541071	
## P	0.1699549931	0.95253658	0.77101391	0.93725640	0.683227214	
## plus_minus	0.0092759827	0.43419903	0.40270019	0.49063225	0.122621012	
## PIM	1.0000000000	0.09915262	0.18623755	0.18749468	0.074706109	
## PperGP	0.0991526176	1.00000000	0.71810767	0.89306021	0.651392559	
## EVG	0.1862375530	0.71810767	1.00000000	0.84291157	0.397107445	
## EVP	0.1874946773	0.89306021	0.84291157	1.00000000	0.497566982	
## PPG	0.0747061093	0.65139256	0.39710744	0.49756698	1.0000000000	
## PPP	0.0988448071	0.81325894	0.47645129	0.62037529	0.808313802	
## SHG	-0.0266492540	0.06520372	0.05813610	0.05440679	0.001476032	
## SHP	0.0464283757	0.08921445	0.09252494	0.08092964	-0.019857744	
## OTG	0.0006747926	0.30664446	0.35425113	0.32611068	0.213541309	
## GWG	0.0978716537	0.54815737	0.59357441	0.56603193	0.467347969	
## S	0.2396343195	0.72348917	0.72226140	0.77673376	0.579858069	
## TOI	0.1520555116	0.81595070	0.62153922	0.76014044	0.573594934	
##	PPP	SHG	SHP	OTG	GWG	
## Salary	0.63118408	0.063562572	0.08757546	0.2273830499	0.41131968	
## GP	0.36409513	0.133123714	0.16823725	0.1391359582	0.31020592	
## G	0.71032409	0.151338573	0.14683281	0.3497504482	0.64582130	
## A	0.82328113	0.036888187	0.08361223	0.2457119608	0.45052559	
## P	0.85056724	0.091518905	0.11976247	0.3154147174	0.58035617	
## plus_minus	0.16895420	0.111059858	0.15834274	0.2045854584	0.35110797	
## PIM	0.09884481	-0.026649254	0.04642838	0.0006747926	0.09787165	
## PperGP	0.81325894	0.065203718	0.08921445	0.3066444649	0.54815737	
## EVG	0.47645129	0.058136100	0.09252494	0.3542511254	0.59357441	
## EVP	0.62037529	0.054406786	0.08092964	0.3261106828	0.56603193	
## PPG	0.80831380	0.001476032	-0.01985774	0.2135413094	0.46734797	
## PPP	1.00000000	0.010659768	0.01006036	0.2276268759	0.46761493	
## SHG	0.01065977	1.0000000000	0.82193297	-0.0103139446	0.07365052	
## SHP	0.01006036	0.821932969	1.00000000	0.0036239001	0.02027842	
## OTG	0.22762688	-0.010313945	0.00362390	1.0000000000	0.48891424	
## GWG	0.46761493	0.073650522	0.02027842	0.4889142439	1.0000000000	
## S	0.66602667	0.113038315	0.11807334	0.3406281889	0.51768473	

```
## TOI      0.70903205  0.108284007  0.14575022  0.2644176808 0.45390051
##          S      TOI
## Salary   0.5801027 0.6703051
## GP       0.6549627 0.5355405
## G        0.7934833 0.7179934
## A        0.7074770 0.7746784
## P        0.8127386 0.8225208
## plus_minus 0.2021904 0.2149741
## PIM      0.2396343 0.1520555
## PperGP   0.7234892 0.8159507
## EVG      0.7222614 0.6215392
## EVP      0.7767338 0.7601404
## PPG      0.5798581 0.5735949
## PPP      0.6660267 0.7090321
## SHG      0.1130383 0.1082840
## SHP      0.1180733 0.1457502
## OTG      0.3406282 0.2644177
## GWG      0.5176847 0.4539005
## S        1.0000000 0.7770869
## TOI      0.7770869 1.0000000
```

```
corrplot(cor(Forwards_num_data), method = "circle", type = "upper", title = "
Forwards")
```



```
# For defensemen
sapply(Def, is.numeric)
```

```
## i..Player      Team      Salary      S.C      Pos      GP
G
##      FALSE      FALSE      TRUE      FALSE      FALSE      TRUE      TR
UE
##      A      P plus_minus      PIM      PperGP      EVG      E
VP
##      TRUE      TRUE      TRUE      TRUE      TRUE      TRUE      TR
UE
##      PPG      PPP      SHG      SHP      OTG      GWG
S
##      TRUE      TRUE      TRUE      TRUE      TRUE      TRUE      TR
UE
##      TOI
##      TRUE
```

```
Def_num_data <- Def[, sapply(Def, is.numeric)]
cor(Def_num_data, use = "complete.obs", method = "pearson")
```


```
##      Salary      GP      G      A      P      plus_minu
s
## Salary      1.00000000 0.33035269 0.4387764 0.4623040 0.5020896 0.04765473
1
## GP      0.33035269 1.00000000 0.3843629 0.5926371 0.5933326 0.07805371
7
## G      0.43877638 0.38436294 1.0000000 0.5434578 0.7269852 0.18829865
3
## A      0.46230400 0.59263707 0.5434578 1.0000000 0.9714875 0.15090415
8
## P      0.50208965 0.59333262 0.7269852 0.9714875 1.0000000 0.17662171
0
## plus_minus 0.04765473 0.07805372 0.1882987 0.1509042 0.1766217 1.00000000
0
## PIM      0.18768978 0.52503350 0.2243964 0.2531786 0.2704769 -0.01156341
1
## PperGP      0.53892738 0.37613683 0.6806634 0.8837416 0.9151414 0.18670947
9
## EVG      0.33423738 0.36164743 0.9171186 0.4542442 0.6306000 0.24667509
7
## EVP      0.44387831 0.65591805 0.6421331 0.8830397 0.9036846 0.26489883
0
## PPG      0.40443912 0.24210271 0.7048778 0.4476925 0.5652954 -0.00692995
4
## PPP      0.42472169 0.34186677 0.6235477 0.8138997 0.8418794 0.00362728
3
## SHG      0.18400309 0.11760204 0.1232040 0.1349607 0.1451947 0.08075430
0
## SHP      0.16712284 0.14669964 0.1697434 0.1160291 0.1428534 0.13877775
4
## OTG      0.12318907 0.08805761 0.1989225 0.2163868 0.2331867 0.18898242
3
```



## GWG 7	0.24189822	0.24618945	0.6073766	0.2541496	0.3794403	0.16818053
## S 2	0.51731854	0.74335396	0.6388341	0.7823879	0.8204205	0.09919940
## TOI 8	0.58633036	0.59409893	0.5513459	0.7074666	0.7344251	0.10259192
## PPP	PIM	PperGP	EVG	EVP	PPG	
## Salary 21686	0.18768978	0.5389274	0.33423738	0.44387831	0.404439124	0.4247
## GP 66771	0.52503350	0.3761368	0.36164743	0.65591805	0.242102713	0.3418
## G 47654	0.22439642	0.6806634	0.91711857	0.64213313	0.704877794	0.6235
## A 99747	0.25317856	0.8837416	0.45424418	0.88303966	0.447692531	0.8138
## P 79393	0.27047693	0.9151414	0.63060000	0.90368464	0.565295396	0.8418
## plus_minus 27283	-0.01156341	0.1867095	0.24667510	0.26489883	-0.006929954	0.0036
## PIM 38863	1.00000000	0.1488959	0.21420917	0.30246589	0.117079106	0.1425
## PperGP 49261	0.14889586	1.00000000	0.59359570	0.80961509	0.528272134	0.7927
## EVG 68385	0.21420917	0.5935957	1.00000000	0.63268491	0.373130938	0.4451
## EVP 23069	0.30246589	0.8096151	0.63268491	1.00000000	0.370918958	0.5343
## PPG 39576	0.11707911	0.5282721	0.37313094	0.37091896	1.000000000	0.6560
## PPP 00000	0.14253886	0.7927493	0.44516838	0.53432307	0.656039576	1.0000
## SHG 72927	0.17177863	0.1111783	0.09136408	0.10812851	-0.055896301	0.1018
## SHP 66063	0.14897732	0.1270469	0.15017056	0.08677591	0.046104325	0.0668
## OTG 69047	0.01106507	0.2113783	0.20230596	0.23047938	0.113197510	0.1792
## GWG 07616	0.23192432	0.3245068	0.56744527	0.32408227	0.392595209	0.3265
## S 57373	0.42019891	0.6706758	0.54338204	0.77906076	0.500394376	0.6378
## TOI 86137	0.25957973	0.6280183	0.46909289	0.70720705	0.439931956	0.5538
## OI	SHG	SHP	OTG	GWG	S	T
## Salary 04	0.18400309	0.16712284	0.12318907	0.2418982	0.5173185	0.58633
## GP 89	0.11760204	0.14669964	0.08805761	0.2461894	0.7433540	0.59409

## G	0.12320400	0.16974337	0.19892251	0.6073766	0.6388341	0.55134
59						
## A	0.13496065	0.11602914	0.21638685	0.2541496	0.7823879	0.70746
66						
## P	0.14519467	0.14285342	0.23318670	0.3794403	0.8204205	0.73442
51						
## plus_minus	0.08075430	0.13877775	0.18898242	0.1681805	0.0991994	0.10259
19						
## PIM	0.17177863	0.14897732	0.01106507	0.2319243	0.4201989	0.25957
97						
## PperGP	0.11117827	0.12704689	0.21137834	0.3245068	0.6706758	0.62801
83						
## EVG	0.09136408	0.15017056	0.20230596	0.5674453	0.5433820	0.46909
29						
## EVP	0.10812851	0.08677591	0.23047938	0.3240823	0.7790608	0.70720
71						
## PPG	-0.05589630	0.04610432	0.11319751	0.3925952	0.5003944	0.43993
20						
## PPP	0.10187293	0.06686606	0.17926905	0.3265076	0.6378574	0.55388
61						
## SHG	1.00000000	0.47462149	-0.01727737	0.1695899	0.2101078	0.13611
10						
## SHP	0.47462149	1.00000000	-0.03871344	0.2069646	0.1665362	0.19270
24						
## OTG	-0.01727737	-0.03871344	1.00000000	0.3021632	0.1408493	0.16488
72						
## GWG	0.16958994	0.20696460	0.30216325	1.00000000	0.3917126	0.30206
38						
## S	0.21010783	0.16653623	0.14084929	0.3917126	1.00000000	0.78172
29						
## TOI	0.13611104	0.19270238	0.16488716	0.3020638	0.7817229	1.00000
00						

```
corrplot(cor(Def_num_data), method = "circle", type = "upper", title = "Defen
semen")
```



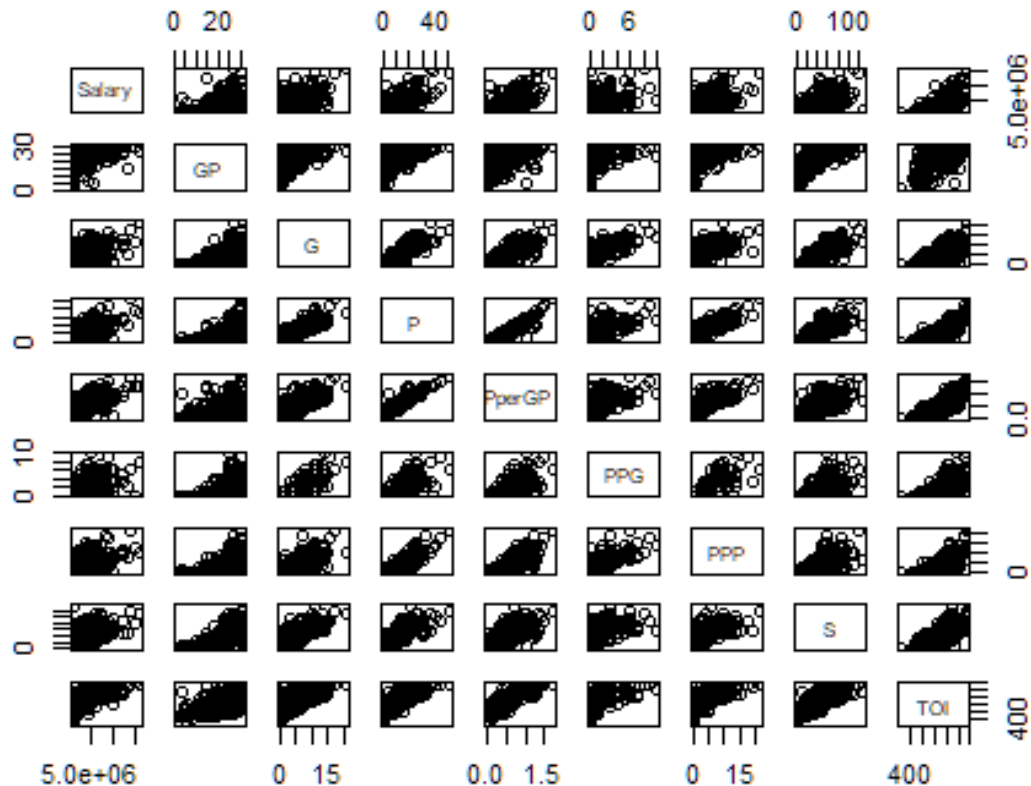
*# Focus on Forwards and remove predictors that show no strong correlation and name this dataframe FWD*

```
FWD <- Forwards_num_data %>% select(Salary, GP, G, P, PperGP, PPG, PPP, S, TOI)
```

```
head(FWD)
```

```
##      Salary GP  G  P PperGP PPG PPP  S  TOI
## 1 15900000 27 21 36   1.33   8  10 112 1323
## 2 15000000 30 11 39   1.30   0  13  82 1358
## 3 14000000 30 17 52   1.73   6  22 119 1349
## 4 13000000 15  5 19   1.27   2   7  45 1229
## 5 12000000 26 13 28   1.08   5  10  87 1210
## 7 12000000 30  9 26   0.87   5   9  83 1090
```

```
pairs(FWD)
```



This gives us an idea which predictors are influential, and which ones are not helpful in predicting the salary. Note that points is autocorrelated with goals and assists.

I will only consider forwards from now on, because the predictors seem better suited to evaluate their value.

*# Create a linear model with promising predictors*

```
FWDModel = lm(Salary~ GP+P+PPP+S+TOI, data=FWD)
```

```
summary(FWDModel)
```

```
##
```

```
## Call:
```

```
## lm(formula = Salary ~ GP + P + PPP + S + TOI, data = FWD)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -4440804 -1056066 -198149   790465  8511183
```

```
##
```

```
## Coefficients:
```


```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) -1457163.3   652315.0  -2.234  0.02607 *
```

```
## GP          -67450.2    21208.2  -3.180  0.00159 **
```

```
## P          88636.9    29537.9    3.001    0.00287 **
## PPP        98434.8    49515.6    1.988    0.04753 *
## S          8796.5     7873.5    1.117    0.26460
## TOI        4779.8      874.7    5.465    8.37e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1872000 on 383 degrees of freedom
## Multiple R-squared:  0.52, Adjusted R-squared:  0.5137
## F-statistic: 82.97 on 5 and 383 DF,  p-value: < 2.2e-16

# Plot residuals - Salaries are not normally distributed
plot(FWDModel$residuals)
abline(h=0, lty=2)
```



Looking at the residuals, this does not look like a normal distribution, as we suspected. I will not pursue linear regression. Log-transforming the salary did not help either, data not shown. Let's try other approaches.

First, I am creating a training and a test set.

*# I tried log-transforming salaries and it doesn't work either. Hence, I will abandon linear regression and will instead focus on other techniques.*

```
options(java.parameters = "-Xmx4g")
```

*# I will use RMSE as my metric, and will define it as follows:*

```
rmse = function(actual, predicted) {  
  sqrt(mean((actual - predicted) ^ 2))  
}
```

*# Then we split the dataset, 80:20 in this case, i.e., 80% of the data will go to the training set and 20% will go to the test set.*

```
FWD1 = sort(sample(nrow(FWD), nrow(FWD)*.8))
```

*#creating training data set by selecting the output row values*

```
train<-FWD[FWD1,]
```

*#creating test data set by not selecting the output row values*

```
test<-FWD[-FWD1,]
```

```
head(train)
```

```
##      Salary GP  G  P PperGP PPG PPP  S  TOI  
## 1 15900000 27 21 36  1.33  8  10 112 1323  
## 3 14000000 30 17 52  1.73  6  22 119 1349  
## 4 13000000 15  5 19  1.27  2   7  45 1229  
## 5 12000000 26 13 28  1.08  5  10  87 1210  
## 7 12000000 30  9 26  0.87  5   9  83 1090  
## 8 12000000 27 10 30  1.11  3  11  49 1132
```

```
nrow(train)
```

```
## [1] 311
```

```
head(test)
```

```
##      Salary GP  G  P PperGP PPG PPP  S  TOI  
## 2  15000000 30 11 39  1.30  0  13  82 1358  
## 13 10000000 26 14 19  0.73  6   6  55 1034  
## 20  9500000 24 12 25  1.04  6  10  63 1109  
## 32  8000000 28 12 24  0.86  3   6  69 1179  
## 37  8000000 25  3 12  0.48  0   2  54  957  
## 45  7000000 25 12 28  1.12  7   8  60 1004
```

```
nrow(test)
```

```
## [1] 78
```

Random Forest.

*# Random Forest*

*# We set up cross-validation 5 fold and create a grid of mtry values (Here, t*

```

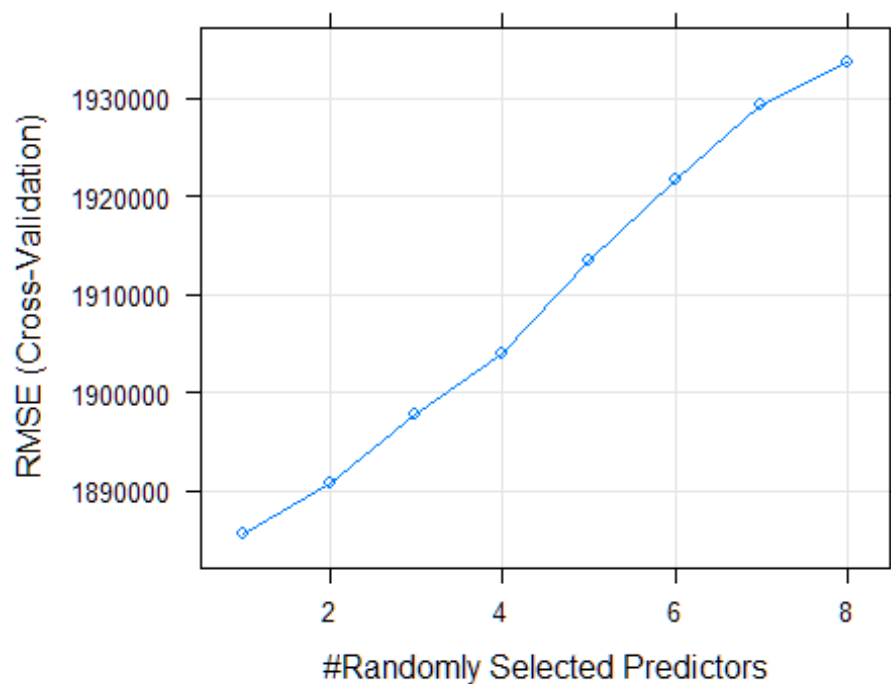
rying all possible values.)
cv_5 = trainControl(method = "cv", number = 5)
rf_grid = expand.grid(mtry = 1:8)

rf_fit = train(Salary ~ ., data = train,
               method = "rf",
               trControl = cv_5,
               tuneGrid = rf_grid)
rf_fit$bestTune

##   mtry
## 1    1

plot(rf_fit)

```



```

rmse(predict(rf_fit, test), test$Salary)

## [1] 1804505

# The resulting test RMSE with mtry = 2 is 1791422

```

Extremely Randomized Trees. Sorry won't run with this PC because my Java version is incompatible. Not the best approach anyways. Code works if the software environment is correct. RMSE: 2014886.


Next, Generalized Boosted Regression Modeling, gbm.

```
# Generalized Boosted Regression Modeling, gbm
gbm_fit = train(Salary ~ ., data = train,
                method = "gbm",
                trControl = cv_5,
                verbose = FALSE,
                tuneLength = 10)

gbm_fit$bestTune

##      n.trees interaction.depth shrinkage n.minobsinnode
## 1         50                1        0.1              10

plot(gbm_fit)
```



```
rmse(predict(gbm_fit, test), test$Salary)

## [1] 1792469
```

Extreme Gradient Boosting, xgboost. Tons of warning messages while processing.


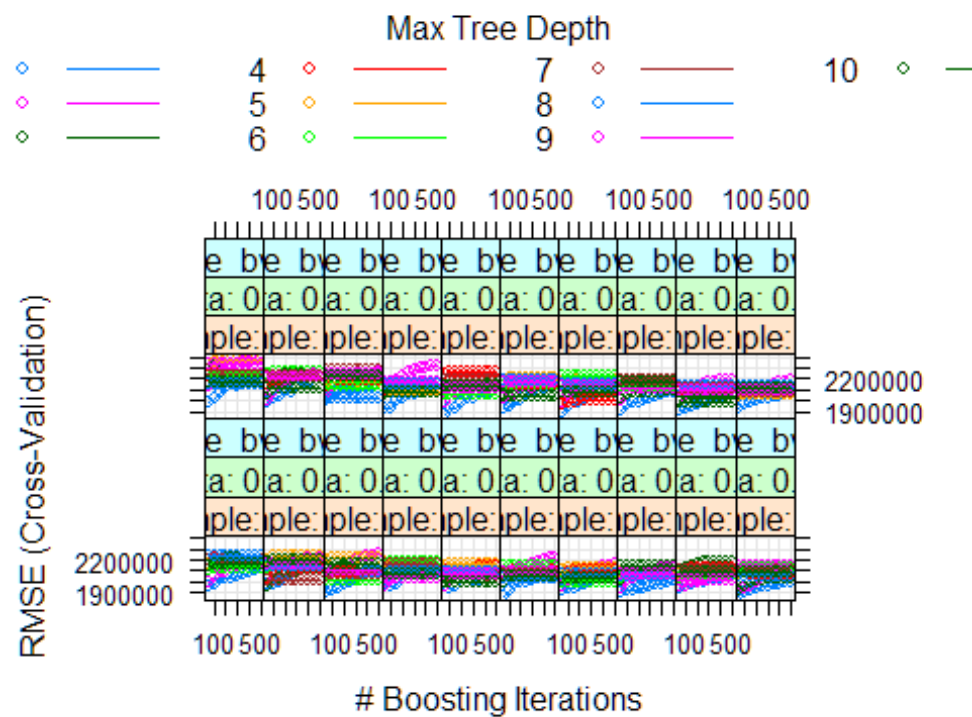
```
# Extreme Gradient Boosting, xgboost. I get tons of warning messages.
xgb_fit = train(Salary ~ ., data = train,
                method = "xgbTree",
                trControl = cv_5,
```



```
verbose = FALSE,  
tuneLength = 10,  
numThreads = 8)
```

-warning messages deleted-

```
plot(xgb_fit)
```



```
rmse(predict(xgb_fit, test), test$Salary)
```

```
## [1] 1796426
```

Finally, I wanted to have a comparison with h2o models. Turns out these aren't much better than the models described above.

```
# Finally, let's see how h2o models compare
```

```
h2o.init(nthreads = -1)
```

```
## Connection successful!
```

```
##
```

```
## R is connected to the H2O cluster:
```

```
## H2O cluster uptime: 1 days 2 hours
```

```
## H2O cluster timezone: Europe/Berlin
```

```
## H2O data parsing timezone: UTC
```

```
## H2O cluster version: 3.32.0.5
```

```
## H2O cluster version age: 2 days
```

```
## H2O cluster name: H2O_started_from_R_ughac_dkv520
```

```
## H2O cluster total nodes: 1
```

```
## H2O cluster total memory: 3.33 GB
```

```
## H2O cluster total cores: 12
```

```
## H2O cluster allowed cores: 12
```

```
## H2O cluster healthy: TRUE
```

```
## H2O Connection ip: localhost
```

```
## H2O Connection port: 54321
```

```
## H2O Connection proxy: NA
```

```
## H2O Internal Security: FALSE
```

```
## H2O API Extensions: Amazon S3, Algos, AutoML, Core V3, TargetEncoder, Core V4
```

```
## R Version: R version 4.0.4 (2021-02-15)
```

```
Fwd_concise = read.csv("Fwd_concise.csv", header = T)
```

```
head(Fwd_concise)
```

```
##           Player Team  Salary S.C Pos GP  G  A  P plus_minus PIM PperG
P EVG
## 1  Auston Matthews  TOR 15900000  L  C 27 21 15 36           8   6   1.3
3 13
## 2  Connor McDavid  EDM 14000000  L  C 30 17 35 52          10  14   1.7
3 11
## 3   John Tavares   TOR 12000000  L  C 30  9 17 26          11   6   0.8
7  4
## 4 Nicklas Backstrom WSH 12000000  L  C 27 10 20 30          -1   4   1.1
1  7
## 5  Sebastian Aho   CAR 10570000  L  C 27 10 15 25           7  12   0.9
3  6
## 6   Jack Eichel   BUF 10000000  R  C 21  2 16 18          -9   6   0.8
6  1
##   EVP PPG PPP SHG SHP OTG GWG  S  TOI
## 1  26   8  10   0   0   1   7 112 1323
## 2  30   6  22   0   0   1   4 119 1349
## 3  17   5   9   0   0   0   1  83 1090
## 4  19   3  11   0   0   0   0  49 1132
```

```
## 5 16 3 8 1 1 0 2 66 1160
## 6 9 1 9 0 0 0 1 61 1230

names(Fwd_concise) <- c("Player", "Salary", "GP", "G", "P", "PperGP", "PPG",
"PPP", "S", "TOI")
head(names)

##
## 1 .Primitive("names")

FWD_h2o <- h2o.importFile("FWD_concise.csv")

## |
| | 0%
|
|=====| 100%

# Create the training dataset and test dataset (80% and 20%)
partitions <- h2o.splitFrame(data = as.h2o(FWD_h2o),
                             ratios = c(0.8),
                             seed = 1)

data_train_h2o <- h2o.assign(data = partitions[[1]], key = "data_train_H2O"
)
data_test_h2o <- h2o.assign(data = partitions[[2]], key = "data_test_H2O")
y1 <- "Salary"
x1 <- setdiff(names(data_train_h2o), y1)

# Applies the H2O AutoML Machine Learning Platform
aml <- h2o.automl(x = x1, y = y1,
                 training_frame = data_train_h2o,
                 validation_frame = data_test_h2o,
                 stopping_metric = "RMSE",
                 seed = 1,
                 sort_metric = "RMSE")

## |
| | 0%
## 22:49:47.685: User specified a validation frame with cross-validation still
enabled. Please note that the models will still be validated using cross-validation
only, the validation frame will be used to provide purely informative validation
metrics on the trained models.
## 22:49:47.685: AutoML: XGBoost is not available; skipping it. |
|=====| 12%
|
|=====| 21%
|
|=====| 25%
|
|=====| 33%
|
```

=====	62%
=====	63%
=====	64%
=====	65%
=====	65%
=====	66%
=====	71%
=====	72%
=====	72%
=====	73%
=====	74%
=====	75%
=====	79%
=====	80%
=====	81%
=====	82%
=====	82%
=====	83%
=====	84%
=====	85%
=====	85%
=====	86%
=====	87%
=====	88%
=====	92%

```
lb <- aml@leaderboard
print(lb, n = nrow(lb))
```

		model_id	rmse
## 1		XRT_1_AutoML_20210318_224947	1923438
## 2	StackedEnsemble_BestOfFamily_AutoML_20210318_224947		1925153
## 3		DRF_1_AutoML_20210318_224947	1937265
## 4		GBM_1_AutoML_20210318_224947	1976159
## 5	GBM_grid__1_AutoML_20210318_224947_model_5		1979149
## 6		GBM_4_AutoML_20210318_224947	1989975
## 7	GBM_grid__1_AutoML_20210318_224947_model_7		1992915
## 8		GBM_2_AutoML_20210318_224947	1995185
## 9		GBM_3_AutoML_20210318_224947	1996622
## 10	GBM_grid__1_AutoML_20210318_224947_model_6		2009496
## 11	GBM_grid__1_AutoML_20210318_224947_model_2		2015722
## 12	GBM_grid__1_AutoML_20210318_224947_model_4		2022996
## 13	GBM_grid__1_AutoML_20210318_224947_model_1		2033074
## 14	DeepLearning_grid__1_AutoML_20210318_224947_model_4		2040069
## 15	DeepLearning_grid__2_AutoML_20210318_224947_model_6		2086859
## 16	DeepLearning_grid__1_AutoML_20210318_224947_model_5		2094879
## 17	DeepLearning_grid__2_AutoML_20210318_224947_model_4		2104743
## 18	DeepLearning_grid__3_AutoML_20210318_224947_model_6		2107548
## 19	DeepLearning_grid__3_AutoML_20210318_224947_model_4		2115614
## 20	DeepLearning_grid__1_AutoML_20210318_224947_model_3		2119638
## 21	DeepLearning_grid__3_AutoML_20210318_224947_model_2		2123925
## 22	DeepLearning_grid__1_AutoML_20210318_224947_model_7		2130639
## 23	DeepLearning_grid__2_AutoML_20210318_224947_model_5		2131762
## 24		GBM_5_AutoML_20210318_224947	2142509
## 25	DeepLearning_grid__1_AutoML_20210318_224947_model_1		2144648
## 26	DeepLearning_grid__1_AutoML_20210318_224947_model_6		2146826
## 27	DeepLearning_grid__1_AutoML_20210318_224947_model_2		2148779
## 28	DeepLearning_grid__2_AutoML_20210318_224947_model_2		2148787
## 29	GBM_grid__1_AutoML_20210318_224947_model_3		2151576
## 30	DeepLearning_grid__2_AutoML_20210318_224947_model_7		2216577
## 31	DeepLearning_grid__3_AutoML_20210318_224947_model_1		2219481
## 32	DeepLearning_grid__2_AutoML_20210318_224947_model_1		2221836
## 33	DeepLearning_grid__3_AutoML_20210318_224947_model_9		2222617
## 34	DeepLearning_grid__3_AutoML_20210318_224947_model_14		2236299
## 35	DeepLearning_grid__3_AutoML_20210318_224947_model_11		2240925
## 36	DeepLearning_grid__3_AutoML_20210318_224947_model_12		2254824
## 37	DeepLearning_grid__3_AutoML_20210318_224947_model_3		2270675
## 38	DeepLearning_grid__3_AutoML_20210318_224947_model_7		2279337
## 39	DeepLearning_grid__3_AutoML_20210318_224947_model_5		2280686
## 40	DeepLearning_grid__3_AutoML_20210318_224947_model_8		2330969
## 41	DeepLearning_1_AutoML_20210318_224947		2331625
## 42	DeepLearning_grid__2_AutoML_20210318_224947_model_3		2363551
## 43	DeepLearning_grid__3_AutoML_20210318_224947_model_10		2470764

```

## 44      StackedEnsemble_AllModels_AutoML_20210318_224947 2768827
## 45      GLM_1_AutoML_20210318_224947 2770200
## 46 DeepLearning_grid__3_AutoML_20210318_224947_model_13 3183626
##      mean_residual_deviance      mse      mae      rmsle
## 1      3.699614e+12 3.699614e+12 1379985 0.6361620
## 2      3.706213e+12 3.706213e+12 1364086 0.6238121
## 3      3.752996e+12 3.752996e+12 1331328 0.6168050
## 4      3.905205e+12 3.905205e+12 1379592 0.6477449
## 5      3.917031e+12 3.917031e+12 1391064 0.6316241
## 6      3.960001e+12 3.960001e+12 1395110 0.6287476
## 7      3.971711e+12 3.971711e+12 1394647 0.6307653
## 8      3.980764e+12 3.980764e+12 1408327 0.6232471
## 9      3.986501e+12 3.986501e+12 1406207 0.6324638
## 10     4.038073e+12 4.038073e+12 1408592 0.6287941
## 11     4.063136e+12 4.063136e+12 1403251 0.6359773
## 12     4.092515e+12 4.092515e+12 1445739 0.6642061
## 13     4.133389e+12 4.133389e+12 1435456 0.6381711
## 14     4.161883e+12 4.161883e+12 1492363 0.7007795
## 15     4.354979e+12 4.354979e+12 1471578 0.6999430
## 16     4.388519e+12 4.388519e+12 1577195      NaN
## 17     4.429944e+12 4.429944e+12 1483405      NaN
## 18     4.441759e+12 4.441759e+12 1528982 0.7242408
## 19     4.475823e+12 4.475823e+12 1561941 0.7391848
## 20     4.492863e+12 4.492863e+12 1537420      NaN
## 21     4.511056e+12 4.511056e+12 1533009 0.6969662
## 22     4.539621e+12 4.539621e+12 1556328      NaN
## 23     4.544411e+12 4.544411e+12 1581337 0.7381308
## 24     4.590344e+12 4.590344e+12 1485106 0.6669821
## 25     4.599517e+12 4.599517e+12 1548487      NaN
## 26     4.608863e+12 4.608863e+12 1546240      NaN
## 27     4.617250e+12 4.617250e+12 1606584      NaN
## 28     4.617284e+12 4.617284e+12 1617629 0.7409023
## 29     4.629278e+12 4.629278e+12 1495947 0.6704016
## 30     4.913213e+12 4.913213e+12 1670437 0.7569916
## 31     4.926096e+12 4.926096e+12 1684527 0.7682585
## 32     4.936554e+12 4.936554e+12 1650548 0.7689965
## 33     4.940024e+12 4.940024e+12 1572785      NaN
## 34     5.001034e+12 5.001034e+12 1648111      NaN
## 35     5.021743e+12 5.021743e+12 1702681 0.7646135
## 36     5.084232e+12 5.084232e+12 1661689 0.7544574
## 37     5.155967e+12 5.155967e+12 1652783 0.7612042
## 38     5.195377e+12 5.195377e+12 1741074 0.7842688
## 39     5.201531e+12 5.201531e+12 1768061 0.7913690
## 40     5.433418e+12 5.433418e+12 1630328 0.7594464
## 41     5.436475e+12 5.436475e+12 1767910      NaN
## 42     5.586372e+12 5.586372e+12 1811536 0.8285262
## 43     6.104674e+12 6.104674e+12 2004179 0.9080212
## 44     7.666403e+12 7.666403e+12 2143457 0.9391688
## 45     7.674006e+12 7.674006e+12 2143914 0.9394052
## 46     1.013548e+13 1.013548e+13 2656635      NaN

```

```

##
## [46 rows x 6 columns]

aml@leader

## Model Details:
## =====
##
## H2ORegressionModel: drf
## Model ID: XRT_1_AutoML_20210318_224947
## Model Summary:
##   number_of_trees number_of_internal_trees model_size_in_bytes min_depth
## 1                35                35                94972            16
##   max_depth mean_depth min_leaves max_leaves mean_leaves
## 1          20   19.17143        146        213   195.00000
##
##
## H2ORegressionMetrics: drf
## ** Reported on training data. **
## ** Metrics reported on Out-Of-Bag training samples **
##
## MSE: 3.978484e+12
## RMSE: 1994614
## MAE: 1441659
## RMSLE: 0.6565352
## Mean Residual Deviance : 3.978484e+12
##
##
## H2ORegressionMetrics: drf
## ** Reported on validation data. **
##
## MSE: 3.144704e+12
## RMSE: 1773331
## MAE: 1239874
## RMSLE: 0.6654103
## Mean Residual Deviance : 3.144704e+12
##
##
## H2ORegressionMetrics: drf
## ** Reported on cross-validation data. **
## ** 5-fold cross-validation on training data (Metrics computed for combined
holdout predictions) **
##
## MSE: 3.699614e+12
## RMSE: 1923438
## MAE: 1379985
## RMSLE: 0.636162
## Mean Residual Deviance : 3.699614e+12
##
##

```



```
## Cross-Validation Metrics Summary:
##
##          mean          sd    cv_1_valid    cv_2_v
alid
## mae          1379842.6    143136.69    1210061.6    15941
62.2
## mean_residual_deviance 3.69990618E12 5.47875029E11 3.02002104E12 4.2887500
3E12
## mse          3.69990618E12 5.47875029E11 3.02002104E12 4.2887500
3E12
## r2          0.5041692    0.103992954    0.6245669    0.535
1587
## residual_deviance    3.69990618E12 5.47875029E11 3.02002104E12 4.2887500
3E12
## rmse          1919185.6    144191.27    1737820.8    20709
29.8
## rmsle          0.63391995    0.059637815    0.5621784    0.7029
0357
##          cv_3_valid    cv_4_valid    cv_5_valid
## mae          1428679.9    1348144.8    1318164.9
## mean_residual_deviance 3.90807054E12 4.05456185E12 3.22812799E12
## mse          3.90807054E12 4.05456185E12 3.22812799E12
## r2          0.5564315    0.35579178    0.4488973
## residual_deviance    3.90807054E12 4.05456185E12 3.22812799E12
## rmse          1976884.0    2013594.2    1796699.2
## rmsle          0.5824985    0.650782    0.67123735

# test prediction of the leader model
pred <- h2o.predict(aml, data_test_h2o)

## |
|
|
|=====| 100%

# retrieve the leaderboard
lb <- h2o.get_leaderboard(object = aml, extra_columns = 'ALL')
lb

##          model_id    rmse
## 1          XRT_1_AutoML_20210318_224947 1923438
## 2 StackedEnsemble_BestOfFamily_AutoML_20210318_224947 1925153
## 3          DRF_1_AutoML_20210318_224947 1937265
## 4          GBM_1_AutoML_20210318_224947 1976159
## 5          GBM_grid__1_AutoML_20210318_224947_model_5 1979149
## 6          GBM_4_AutoML_20210318_224947 1989975
## mean_residual_deviance    mse    mae    rmsle training_time_ms
## 1          3.699614e+12 3.699614e+12 1379985 0.6361620    143
## 2          3.706213e+12 3.706213e+12 1364086 0.6238121    127
## 3          3.752996e+12 3.752996e+12 1331328 0.6168050    89
## 4          3.905205e+12 3.905205e+12 1379592 0.6477449    33
## 5          3.917031e+12 3.917031e+12 1391064 0.6316241    14
```

```
## 6          3.960001e+12 3.960001e+12 1395110 0.6287476          35
##  predict_time_per_row_ms          algo
## 1          0.041775          DRF
## 2          0.042058 StackedEnsemble
## 3          0.007266          DRF
## 4          0.004837          GBM
## 5          0.003211          GBM
## 6          0.004408          GBM
##
## [46 rows x 9 columns]
```

### Conclusions:

There are several methods that produces similar RMSE values, including Random Forest, gbm, xgboost, and h2o models. I could not create a linear regression model. It is inherently difficult to predict salaries in professional sports. There are many reasons why a closer fit may not be achievable, for instance the time when a contract was signed.